

HR Analysis - Employee Churn

After much struggle with my original dataset and ideas, I had to fall back to a contingency plan, still focusing on my main idea of churn, but instead for employee turnover and retention, rather than company contractual churn.

My original idea was to use a real life dataset from my work but I came across two challenges, I was not given ideal data for the challenge of predicting churn, and I ran security risks if I did not conceal the data.

```
In [1]: # import initial needed libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: # read in our data

df = pd.read_csv("HR_Dataset.csv")
df.head()
```

Out[2]:

satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	Work_accident	left	promotion_last_5years
0.38	0.53	2	157	3	0	1	0
0.80	0.86	5	262	6	0	1	0
0.11	0.88	7	272	4	0	1	0
0.72	0.87	5	223	5	0	1	0
0.37	0.52	2	159	3	0	1	0

```
In [3]: df.columns
```

```
Out[3]: Index(['satisfaction_level', 'last_evaluation', 'number_project',
               'average_monthly_hours', 'time_spend_company', 'Work_accident', 'left',
               'promotion_last_5years', 'Departments ', 'salary'],
              dtype='object')
```

```
In [4]: # need to change 'Departments ' to 'Departments'

df.rename(columns = {'Departments ':'Departments'}, inplace = True)
```

```
In [5]: df.shape
```

```
Out[5]: (14999, 10)
```

```
In [6]: # use df.info() to get some information on our data set
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14999 entries, 0 to 14998
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   satisfaction_level      14999 non-null  float64
1   last_evaluation         14999 non-null  float64
2   number_project          14999 non-null  int64
3   average_monthly_hours  14999 non-null  int64
4   time_spend_company      14999 non-null  int64
5   Work_accident           14999 non-null  int64
6   left                    14999 non-null  int64
7   promotion_last_5years   14999 non-null  int64
8   Departments             14999 non-null  object
9   salary                  14999 non-null  object
dtypes: float64(2), int64(6), object(2)
memory usage: 1.1+ MB
```

```
In [7]: # Departments and Salary are both object: categorical variables that will
# need to be transformed before moving on to prediction
# Lets take a look at the values within these columns
```

```
print(df.salary.unique())
```

```
['low' 'medium' 'high']
```

```
In [8]: print(df.Departments.unique())

['sales' 'accounting' 'hr' 'technical' 'support' 'management' 'IT'
 'product_mng' 'marketing' 'RandD']
```

Ordinal variables have two or more categories which can be ranked and ordered

Nominal variables have two or more categories which do not have an intrinsic order

Salary: Ordinal

Department: Nominal

The dataset I chose has no missing values, it is best practice in HR Departments to have all personal info on employees on file, but the process in which a the data is collected and stored can have an impact on its accuracy.

```
In [9]: # start transforming categorical variables to numeric
# in the case of ordinal variables we can convert them into a respective numeric value
# encoding the salary column
# first changing the data type to categorical
df.salary = df.salary.astype('category')
```

```
In [10]: # the ordinal values are low, medium, and high in this order
df.salary = df.salary.cat.reorder_categories(['low', 'medium', 'high'])
```

```
In [11]: # encode the values to make integer
df.salary = df.salary.cat.codes
```



```
In [15]: # need to drop the old department column

df = df.drop("Departments", axis=1)

# join the new dataframe df1 to the dataset

df = df.join(df1)
```

```
In [16]: df.head()
```

```
Out[16]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spend_company	Work_accident	left	promotion_last_5y
0	0.38	0.53	2	157	3	0	1	
1	0.80	0.86	5	262	6	0	1	
2	0.11	0.88	7	272	4	0	1	
3	0.72	0.87	5	223	5	0	1	
4	0.37	0.52	2	159	3	0	1	

```
In [17]: # the variable that tells us whether the employee has churned or not
# is in the column left
# 1 = churn
# 0 = has not churned
# I will need to calculate the turnover rate percentage by counting the times 1 and 0 occurred and divide by total
```

```
In [18]: # total number of observations in the dataset for employees
n_employees = len(df)
```

```
In [19]: # number of employees who have churned or not
print(df.left.value_counts()) # have not churned = 11,428, churn = 3,571


0    11428
1     3571
Name: left, dtype: int64
```

```
In [20]: # percentage of employees churn/not churn
print(df.left.value_counts()/n_employees*100)
```

```
0    76.191746
1    23.808254
Name: left, dtype: float64
```

```
In [21]: # now I need to look at correlations or variables that are in a positive
# or negative linear relationship with the target variable
```

```
corr_matrix = df.corr()
sns.heatmap(corr_matrix)
plt.show()
```



```
In [22]: # in the heat map, you can see that Left has the most filled in,
# or highest negative correlation with satisfaction level
# in otherwords, inversely correlated, increase in satisfaction, means decrease in the probability an employee w
```

```
In [23]: # split the data
target = df.left
features = df.drop("left",axis=1)
from sklearn.model_selection import train_test_split
```

```
In [24]: target_train, target_test, features_train, features_test = train_test_split(target,features,test_size=0.25, rand
```

```
In [26]: # employee churn is a binary classification problem, one algorithm that works
# well with this type of problem is Decision tree model
# decision trees have accurate predictions and can help with understanding the reasons employees churn
# import the decision tree classifier
from sklearn.tree import DecisionTreeClassifier
```

```
In [30]: #instantiate the decision tree classifier and set a random state
# limiting the max depth to 5 levels, to prevent overfitting
clf = DecisionTreeClassifier(max_depth=5, random_state=42) # random state does not affect the model results but
```

```
In [31]: # now that the model is set up, we can use the fit() method to fit our features to the target
# will also use .tree_.node_count to get the number of nodes in the tree
clf.fit(features_train, target_train).tree_.node_count
```

Out[31]: 45

```
In [32]: # to test how good the tree is making predictions we need to calculate
# the accuracy score of the prediction using score()
# because we developed the model on the training set, we calculate accuracy score on the test set
clf.score(features_test,target_test)*100
```

Out[32]: 97.06666666666666

```
In [33]: clf.score(features_train,target_train)*100
```

Out[33]: 97.71535247577563

```
In [ ]: # accuracy score dropped a bit when limiting the growth of the tree,
# but the difference between shows that we reduced overfitting and the model will act more realistically
```

```
In [35]: # giving features_test, data points where I haven't given the model the answers to  
# but I want to get the answers that the model comes up with
```

```
clfpredictions = clf.predict(features_test)  
clfpredictions
```

```
Out[35]: array([0, 0, 0, ..., 0, 0, 0], dtype=int64)
```

```
In [36]: # whole dataset had nearly 15,000 rows, we set aside 3750 for testing  
# when we do predictions, it's going to give us the class number that it is predicting, so 0 or 1, churn or not churn  
# so for the first three as you can see from the array above the model is predicting 0, they did not churn  
# this is a way to see generally how the model is predicting on the new information that it's giving you  
features_test
```

```
Out[36]:
```

	satisfaction_level	last_evaluation	number_project	average_monthly_hours	time_spent_company	Work_accident	promotion_last_5ye
6723	0.65	0.96	5	226	2	1	
6473	0.88	0.80	3	166	2	0	
4679	0.69	0.98	3	214	2	0	
862	0.41	0.47	2	154	3	0	
7286	0.87	0.76	5	254	2	1	
...
10371	0.99	0.37	6	219	6	0	
12541	0.81	0.87	4	254	5	0	
2656	0.67	0.59	3	177	3	1	
6759	0.22	0.57	5	174	6	0	
13564	0.36	0.73	2	111	2	0	

3750 rows × 17 columns




```
In [39]: # so we have the predictions for the test instances but how will we know if the predictions are good or not? we  
# by using our performance metrics
```

```
from sklearn.metrics import accuracy_score  
accuracy_score(target_test, clfpredictions)*100
```

```
Out[39]: 97.06666666666666
```

```
In [42]: from sklearn.metrics import confusion_matrix  
confusion_matrix(target_test, clfpredictions, labels=[0,1])
```

```
# if the class is 0 and is predicted as 0, this happened 2812 times  
# when the class is 1 and predicted as 1 this happened 828 times  
# wrongly classified instances FP=69, and FN 41
```

```
Out[42]: array([[2812,   41],  
               [  69,  828]], dtype=int64)
```

```
In [46]: # our precision score
```

```
from sklearn.metrics import precision_score  
precision_score(target_test, clfpredictions)*100
```

```
Out[46]: 95.2819332566168
```

```
In [47]: # recall score  
# True positives over the sum of True positives and False negatives  
# if our goal is mostly focused on those who are churning, then you probably  
# want to have less false negatives, people who leave in reality but your algorithm is not able to predict it  
# for that reason recall can be useful  
# higher values of recall correspond lower values of false negatives  
# recall score 92%, of correct predictions among 1s(churners)
```

```
from sklearn.metrics import recall_score  
recall_score(target_test, clfpredictions)*100
```

```
Out[47]: 92.3076923076923
```

```
In [ ]: # we want to be able to correctly predict churn, so recall score is our target
# but recall score is not enough, because only targeting one class the 1s we might have low accuracy for our 0s,
# people who did not churn
# so I want to use a measure that is not concentrated on just one class or the other
# AUC score if we want to target both churners and non churners
# Area Under Curve and is basically a compound measure that is maximized when
# both recall and specificity are maximized
# using AUC as a target to maximize, the model will try to correctly classify both
# 1s and 0s keeping an eye on recall and specificity at the same time
#
```

```
In [48]: from sklearn.metrics import roc_auc_score
```

```
In [50]: prediction = clf.predict(features_test)
roc_auc_score(target_test, prediction)*100
```

```
Out[50]: 95.43530426811185
```

```
In [53]: new_model = DecisionTreeClassifier(max_depth=5, class_weight="balanced", random_state=42)
```

```
In [55]: new_model.fit(features_train, target_train).tree_.node_count
```

```
Out[55]: 47
```

```
In [56]: # print the accuracy score of the new model
print(new_model.score(features_test, target_test)*100)
```

```
93.70666666666668
```

```
In [63]: # print the recall score for new model
print(recall_score(target_test, clf.predictions)*100)
```

```
92.3076923076923
```

```
In [ ]: # how do we decide what the max_depth or other parameters should be?
# we can simply do this by trying different values and find the one that provides the
# best predictions
# to find the optimal parameters we can create a grid (referenced sklearn documentation),
# of values we want to test to find the values that will give us the highest accuracy
# with hyperparameter tuning we also want to validate the model on different test components
# using cross validationa
```

```
In [86]: # grid search and cross validation for hyperparameter tuning
from sklearn.model_selection import cross_val_score
print(cross_val_score(clf, features, target, cv=10))

[0.98      0.97133333 0.97266667 0.97333333 0.97133333 0.97266667
 0.97933333 0.96866667 0.96866667 0.96731154]
```

```
In [87]: # generate values for maximum depth
depth = [i for i in range(5,21,1)]
```

```
In [88]: # generate values for minimum sample size
samples = [i for i in range(50,500,50)]
```

```
In [89]: # Create the dictionary with parameters to be checked
parameters = dict(max_depth=depth, min_samples_leaf=samples)
```

```
In [90]: # import grid search from sklearn
from sklearn.model_selection import GridSearchCV
```

```
In [91]: param_search = GridSearchCV(clf, parameters)
```

```
In [92]: param_search.fit(features_train, target_train)
```

```
Out[92]: GridSearchCV(estimator=DecisionTreeClassifier(max_depth=5, random_state=42),  
                      param_grid={'max_depth': [5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15,  
                                                16, 17, 18, 19, 20],  
                                'min_samples_leaf': [50, 100, 150, 200, 250, 300, 350,  
                                                    400, 450]})
```

```
In [93]: print(param_search.best_params_)  
  
{'max_depth': 5, 'min_samples_leaf': 50}
```

```
In [94]: # feature importance  
# what are the important features that drive the decision to leave the company?  
# sklearn can calculate the feature importance  
# in sklearn, importances are scaled up to equal 100%  
# the higher the percentage the higher the importance  
# if you find that a feature is not important at all, you should drop it and run the model without that feature
```

```
In [96]: # Lets find the features in our data  
feature_names = list(features)  
feature_names
```

```
Out[96]: ['satisfaction_level',  
          'last_evaluation',  
          'number_project',  
          'average_monthly_hours',  
          'time_spend_company',  
          'Work_accident',  
          'promotion_last_5years',  
          'salary',  
          'RandD',  
          'accounting',  
          'hr',  
          'management',  
          'marketing',  
          'product_mng',  
          'sales',  
          'support',  
          'technical']
```

```
In [98]: # one of the best things about decision trees is its interpretability
# so you can understand how the decision tree is making its decisions
# so we can see which features are more important than the others


feature_importance = pd.DataFrame(clf.feature_importances_, index = feature_names, columns=["importance"])
feature_importance.sort_values(by="importance", ascending=False)
```

Out[98]:

	importance
satisfaction_level	0.535522
time_spend_company	0.152146
last_evaluation	0.151002
number_project	0.097404
average_monthly_hours	0.062417
technical	0.001510
promotion_last_5years	0.000000
salary	0.000000
Work_accident	0.000000
accounting	0.000000
hr	0.000000
management	0.000000
marketing	0.000000
product_mng	0.000000
sales	0.000000
support	0.000000
RandD	0.000000

```
In [102]: feature_importance.plot(kind='bar')
```

```
Out[102]: <AxesSubplot:>
```



```
In [103]: # now that we can see what features are important  
# we need to select those features that have the importance above 1%  
selected_features = feature_importance[feature_importance.importance>0.01]
```

```
In [104]: # create a List from those features  
selected_list = selected_features.index
```

```
In [105]: # transform both features_train and features_test components
# to include only selected features
features_train_selected = features_train[selected_list]
features_test_selected = features_test[selected_list]
```

```
In [133]: # I now have the most important features, along with the optimal parameters we found earlier
# so now I have to redo the model for predicting churn with these new features and parameters
# (adjusted to get higher accuracy)
best_model = DecisionTreeClassifier(max_depth=5, min_samples_leaf = 100, class_weight="balanced", random_state=4
```

```
In [134]: # fit this new model with only the selected features
best_model.fit(features_train_selected, target_train).tree_.node_count
```

Out[134]: 41

```
In [135]: # make a new prediction based on the important features from the test set
best_prediction = best_model.predict(features_test_selected)
```

```
In [136]: # Lets see the general accuracy of the best_model
print(best_model.score(features_test_selected, target_test)*100)
```

95.62666666666667

```
In [137]: # recall score of best_model
print(recall_score(target_test, best_prediction)*100)
```

92.08472686733556

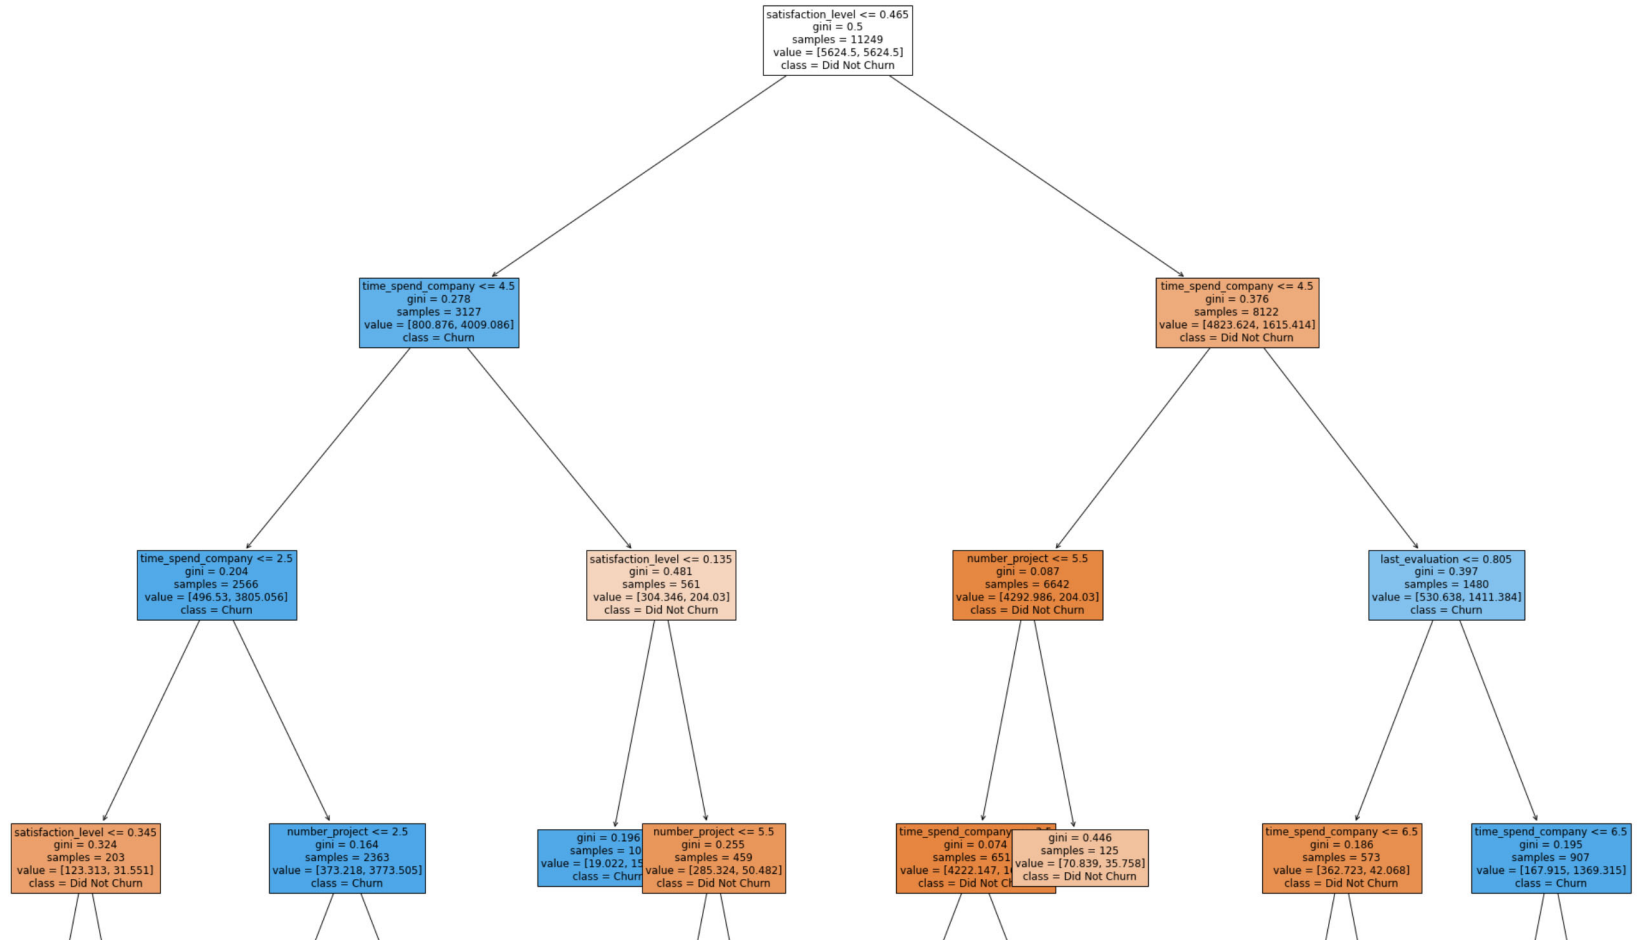
```
In [138]: # print ROC/AUC score of the best_model
print(roc_auc_score(target_test, best_prediction)*100)
```

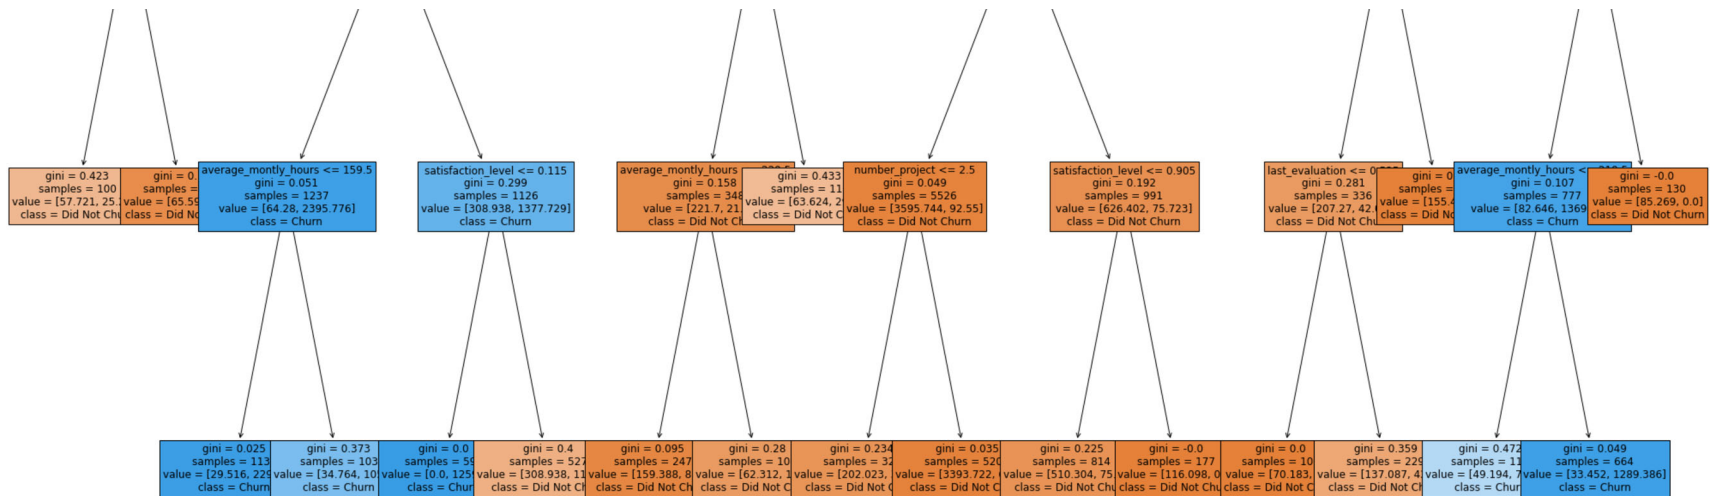
94.4125001318802

```
In [141]: # time to visualize the decision tree
# import from sklearn
from sklearn import tree
from matplotlib import pyplot as plt
```

```
In [149]: fig = plt.figure(figsize=(35,35))
_ = tree.plot_tree(best_model,
                  feature_names=feature_names,
                  class_names={0: 'Did Not Churn', 1: 'Churn'},
                  filled=True,
                  fontsize=12)
```

our most important feature = satisfaction_level
we created a plot of the tree, which is an actual representation of how the tree decides
for example we can see the root node, if satisfaction_level <=0.465, if its Lower, it branches off to the Left
if its higher, it branches off to the right, on to the next decision point and so on.





In [150]: selected_features

Out[150]:

	importance
satisfaction_level	0.535522
last_evaluation	0.151002
number_project	0.097404
average_monthly_hours	0.062417
time_spend_company	0.152146

In []: